

Hospital Chatbot for Color-Coded Clinical Pathways

A Deep Dive into NLP, Transformers, and Multi-Layer Triage

Undergraduate Research Presentation

March 26, 2026

Presentation Outline

1. Introduction & Key NLP Definitions
2. NLP Pipelines
3. The Mathematics of Attention
4. Transformers: Architecture In-Depth
5. Intro to BERT
6. BioBERT & Datasets

7. Layer 1: The NLP Pipeline
8. Layer 2: The Deterministic Model
9. Layer 3: The Bayesian Probabilistic Model
10. Color-Coded Clinical Pathways & Streaming
11. Mathematical Foundations of BERT
12. Conclusion

Introduction to Clinical NLP

- **Definition:** Natural Language Processing (NLP) bridges computer science and linguistics. In healthcare, it transforms highly subjective, unstructured medical text (triage notes) into structured, predictive data.
- **The Core Challenge:** Emergency Department (ED) notes lack strict sequence order. Nurses write rapidly, producing high variance, typos, and abbreviations.
- **Historical Evolution:**
 - *Rule-Based (1950s-80s):* Handcrafted syntax rules. Failed due to linguistic ambiguity.
 - *Statistical (1990s-00s):* Relied on word frequencies (TF-IDF). Ignored context.
 - *Deep Learning (2010s):* RNNs/CNNs learned features automatically.
 - *Transformers (2017-Present) (Vaswani et al., 2017):* Attention-only architecture enabling massive parallelization and contextual understanding.

Key NLP Definitions

To understand the math, we must define the core units of NLP:

- **Tokenization:** The process of breaking raw text into fundamental units (words, sub-words, or characters).
- **Embeddings:** Dense, low-dimensional mathematical vectors (arrays of numbers) representing the semantic meaning of a token.
- **Static vs. Contextual Embeddings:** Older models (Word2Vec) (Mikolov et al., 2013) gave the word "bank" one static vector. Modern models (BERT) (Devlin et al., 2018) generate *contextual* embeddings, where the vector changes dynamically depending on the surrounding sentence.

Traditional vs. Deep Learning Pipelines

Traditional NLP Pipeline

- *Flow:* Tokenization → Part-of-Speech Tagging → Dependency Parsing → Entity Extraction.
- *Limitation:* Highly modular. Errors compound at each sequential stage. Requires manual feature engineering.

Modern Deep Learning Pipeline

- *Flow:* Raw Text → Sub-word Tokenization → Contextual Embedding Layer → Deep Neural Network → Output Probabilities.
- *Advantage:* End-to-end processing. The model learns optimal representations automatically.

Why Attention? (The Clinical Need)

The Problem with Standard Pooling

Traditional neural networks take the mathematical average of all words in a document. This is dangerous in medicine: averaging treats the word "patient" with the same mathematical weight as "gunshot" or "cardiac arrest".

The Solution: Attention models build a dynamic neural layer to mathematically evaluate and score the "importance" of every single word based on its surrounding context.

Additive Attention Formula (Variable Breakdown)

In models like the Deep Attention Model (DAM), attention is calculated via:

$$a(i) = \frac{\exp(s(h_n^{(i)}; \theta))}{\sum_{i=1}^{l_n} \exp(s(h_n^{(i)}; \theta))}$$

Defining the Parameters:

- $a(i)$: The final normalized **Attention Score** (weight) for the i -th word.
- $h_n^{(i)}$: The **Hidden State** (vector representation) of the i -th word.
- θ : The **Trainable Weights** of the neural network learning the attention.
- $s(\dots)$: The **Scoring Function** (a dense neural layer) evaluating the word's importance.
- l_n : The total **Length of the Document** (number of words).
- $\exp(\dots)$: The exponential function, used here within a **Softmax** operation to ensure all word attention scores sum up to exactly 1.0.

Scaled Dot-Product Attention (Transformers)

Transformers upgraded this to "Self-Attention", evaluating how words relate to each other:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$$

Defining the Parameters:

- **Q (Query matrix):** Represents "What is this token looking for?"
- **K (Key matrix):** Represents "What information does this token contain?"
- **V (Value matrix):** The actual semantic content of the token.
- QK^T : The **Dot Product**. This multiplies Queries and Keys to calculate mathematical similarity (alignment) between every pair of words.
- $\sqrt{d_k}$: **Scaling Factor**. d_k is the dimension of the Key vectors. Dividing by its square root prevents the dot products from growing too large and pushing the softmax function into regions with near-zero gradients.

What is a Transformer?

- **Definition:** A deep learning architecture introduced by Google in 2017 ("Attention is All You Need") (Vaswani et al., 2017) that relies *entirely* on self-attention mechanisms, dispensing with recurrent networks (RNNs/LSTMs) completely.
- **The Core Innovation:** Because they do not process text sequentially (left-to-right), Transformers look at an entire sequence simultaneously.
- **The Benefit:** Resolves the vanishing gradient problem for long documents and allows for massive parallel computation on GPUs, leading to models with billions of parameters.

Transformer Architecture Breakdown

A standard Transformer consists of two main blocks:

- **Encoder:** Reads the input text and maps it into a continuous, deeply contextualized mathematical representation.
- **Decoder:** Takes the Encoder's representation and generates an output sequence one token at a time (auto-regressive).

Crucial Sub-components:

- *Positional Encoding:* Because Transformers process everything at once, they have no concept of word order. We mathematically inject sine/cosine waves into the input embeddings to track positions.
- *Multi-Head Attention:* Computes the $Attention(Q, K, V)$ formula multiple times in parallel, allowing the model to simultaneously focus on different relationships (e.g., one head tracks grammar, another tracks clinical symptoms).

Defining BERT

Bidirectional Encoder Representations from Transformers (Devlin et al., 2018).

- **Architecture:** BERT is an **Encoder-Only** Transformer. It drops the Decoder because its goal is to understand text deeply, not to generate new paragraphs.
- **Bidirectional Context:** Unlike previous models that looked left-to-right, BERT's self-attention looks at the entire sentence in both directions simultaneously across all its layers.
- **Transfer Learning Paradigm:** BERT is "pre-trained" on massive general data (Wikipedia), learning the rules of language. We then "fine-tune" it on smaller, specific datasets (like our hospital triage data).

How BERT Learns (Training Objectives)

How do you teach a machine the English language without labeled data?

1 Masked Language Modeling (MLM):

- The model randomly hides 15% of the words in a sentence.
- BERT must use the surrounding context to predict the masked word. (e.g., "The patient presented with a severe _____ pain.")
- Forces deep bidirectional semantic understanding.

2 Next Sentence Prediction (NSP):

- BERT is fed two sentences (A and B). It must predict if B logically follows A.
- Teaches the model relationships spanning across entire documents.

The BioBERT Variant

- **The Motivation:** Standard BERT fails in healthcare because medical literature contains highly specific jargon, chemical compounds, and abbreviations not found in standard English corpora.
- **The Solution (BioBERT):** Researchers took the pre-trained standard BERT and continually pre-trained it for 23 extra days on massive biomedical corpora (Lee et al., 2020).
- **The Datasets:**
 - *PubMed Abstracts:* Over 4.5 billion words of biomedical scientific abstracts.
 - *PubMed Central (PMC):* Over 13.5 billion words from full-text biomedical research articles.
- **Result:** A model that inherently understands the complex semantic relationships between diseases, drugs, and symptoms.

Math of BERT: The Layer Computation

Inside the Transformer, the data passes through L identical layers. For each layer l :

Step 1: Self-Attention & Normalization

$$Z^l = \text{LayerNorm}(H^{l-1} + \text{MultiHead}(H^{l-1}, H^{l-1}, H^{l-1}))$$

- H^{l-1} : The **Input State** from the previous layer.
- $\text{MultiHead}(\dots)$: The attention operation passing H^{l-1} as Query, Key, and Value.
- $\text{LayerNorm}(\dots)$: **Layer Normalization** standardizes the variance of the vectors to stabilize deep network training. Note the residual connection (adding H^{l-1} back in).

Math of BERT: Feed-Forward & Output

Step 2: Non-Linearity

$$H^l = \text{LayerNorm}(Z^l + \text{FFN}(Z^l))$$

- Z^l : The **Intermediate State** outputted by the attention mechanism.
- $\text{FFN}(\dots)$: A two-layer **Feed-Forward Network** with a ReLU or GELU activation function. This applies non-linear transformations to each position identically and independently.
- H^l : The final **Hidden State output** of layer l , passed to the next layer.

Math of BERT: Input Representations

Returning to the core NLP engine, before text enters the Attention layers, it is converted into a base matrix $E(token)$ via element-wise addition:

$$E(token) = E_{word} + E_{position} + E_{segment}$$

Defining the Parameters:

- E_{word} : The base dictionary **Token Embedding** (e.g., the vector for "fever").
- $E_{position}$: The **Positional Embedding**. A learned vector indicating whether the word is 1st, 5th, or 20th in the sentence.
- $E_{segment}$: The **Segment Embedding**. Used during Next Sentence Prediction to tell the model if a word belongs to Sentence A or Sentence B.

Integrating BioBERT into the Hospital Chatbot

Layer 1: The NLP Extraction Process

- ➊ **Input:** Patient inputs raw symptoms into the chatbot interface.
- ➋ **Tokenization:** The system uses WordPiece to break down complex medical terms into recognized sub-words.
- ➌ **Feature Extraction:** BioBERT processes the sequence, using Multi-Head Attention to generate high-dimensional contextual embeddings.
- ➍ **Dimensionality Reduction:** Optionally use Principal Component Analysis (PCA) to compress the massive feature space.
- ➎ **Classification Output:** A final neural classification layer maps the embeddings to predict the initial patient disposition.

Layer 2: The Deterministic Algebraic Layer (TEWS)

After the NLP layer extracts subjective symptoms, we validate the patient's status using objective vital signs inputted by a provider.

The Triage Early Warning Score (TEWS) Equation

$$TEWS = \sum_{k=1}^n w_k \cdot f_k(v_k)$$

Defining the Parameters:

- v_k : The objective **Vital Sign** (e.g., heart rate, temperature).
- $f_k(...)$: A strict **Piecewise Function** that evaluates the vital sign against clinical thresholds to output a specific sub-score.
- w_k : A **Weighting Factor** adjusting the clinical importance of the specific vital sign.

Algorithmic Color Mapping

The deterministic equation entirely removes human calculation error. The final TEWS score strictly dictates the **Color-Coded Pathway**:

- **TEWS 7+** → **Red** (Immediate Resuscitation Required)
- **TEWS 5-6** → **Orange** (Urgent: < 10 minutes)
- **TEWS 3-4** → **Yellow** (Semi-Urgent: < 60 minutes)
- **TEWS 0-2** → **Green** (Divert to Polyclinic / Non-Urgent)

Layer 3: The Bayesian Override

The Clinical Flaw: If an active heart attack patient has temporarily normal resting vitals, the deterministic math scores them as Green (0). What if vitals are missing or misleading?

Bayes' Theorem for Medical Probability

$$P(D|S) = \frac{P(S|D) \cdot P(D)}{P(S)}$$

Defining the Parameters:

- $P(D|S)$ (**Posterior**): The final probability the patient has the severe Disease given the NLP-extracted Symptoms.
- $P(S|D)$ (**Likelihood**): The probability these exact symptoms occur in patients who actually have the disease.
- $P(D)$ (**Prior**): The baseline historical prevalence of the disease.

Executing the Clinical Override

- **The Calculation:** The system evaluates the NLP tokens (e.g., $S = \text{Confusion, Anuria, Cold}$) against historical databases.
- **The Result:** If the Bayesian calculation evaluates to an exceptionally high Posterior Probability (e.g., $P(\text{SepticShock}|S) = 0.92$).
- **The Override:** The system triggers a **Clinical Override**, bypassing the low TEWS score and instantly upgrading the patient from the Green pathway to the Red pathway to prevent under-triage fatalities.

The Concept of Patient Streaming, (SATS Manual, 2012)

Once the final triage color is determined by the 3-Layer architecture, the system enforces strict **Patient Streaming**.

- **Why Stream?** Without streaming, non-urgent (Green) patients are constantly pushed to the back of the queue by incoming emergencies, causing massive delays.
- **Resource Allocation:** Streaming ensures that a patient's physical location matches the resources they require (e.g., separating a patient who needs a chair from a patient who needs a ventilator).
- **The Goal:** High-acuity patients are saved immediately, while low-acuity patients are seen efficiently in an alternative area.

High Acuity Pathways: RED & ORANGE

RED Pathway (Emergency)

- **Target Time:** IMMEDIATE.
- **Management:** Bypasses all queues. Taken directly to the **Resuscitation Room**.
- **Resources:** Requires full physiological monitoring, high-powered interventions, and a full medical team response.

ORANGE Pathway (Very Urgent)

- **Target Time:** < 10 minutes.
- **Management:** Referred to the **Majors** area of the Emergency Department.
- **Resources:** Requires advanced equipment and highly trained personnel to manage imminent life threats.

Lower Acuity & Special Pathways

- **YELLOW Pathway (Urgent):**
 - **Target Time:** < 1 hour.
 - **Management:** Referred to the **Majors** area.
- **GREEN Pathway (Routine/Non-Urgent):**
 - **Target Time:** < 4 hours.
 - **Management:** Streamed to the **Minors** area or a primary care clinic. Requires minimal resources (e.g., a single practitioner and a consultation room).
- **BLUE Pathway (Deceased):**
 - **Target Time:** < 2 hours.
 - **Management:** Dead on arrival. Referred directly to a doctor for legal certification.

Summary: The Three-Layered Architecture

By integrating these three computational models, we transform triage from a bottleneck into a high-speed, mathematically guaranteed safety net:

- **Layer 1 (BioBERT NLP):** Deep contextual understanding of subjective patient narratives.
- **Layer 2 (Deterministic TEWS):** Mathematical security utilizing objective baseline vitals.
- **Layer 3 (Bayesian Override):** Probabilistic hyper-vigilance predicting covert clinical deterioration.

References

- Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2018). BERT: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.
- Lee, J., Yoon, W., Kim, S., Kim, D., Kim, S., So, C. H., & Kang, J. (2020). BioBERT: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4), 1234-1240.
- Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. *Advances in neural information processing systems*, 30.